

Physical Attacks and Countermeasures

2: Introduction

amir.moradi@tu-darmstadt.de

Agenda


- Cryptography and Cryptographic Devices
- Examples of Attacks on Cryptographic Devices
- Physical Attacks
- Basics of Hardware Design

Cryptography

- is the science of writing in secret (too simple?)
 - is the scientific study of techniques for securing digital information, transactions and computations
 - It MUST provide
 - Integrity
 - Confidentiality
 - Authenticity
 - Non-Repudation

Classical Scenario

- Alice and Bob are honest parties and want to communicate
- The communication link between them is not secure
- Eve can listen to the communication



Security


- *“A cryptosystem must not be required to be secret, and it must be able to fall into the hands of the enemy without inconvenience” [Kerckhoffs’s Principle]*
 - The security of cryptosystems is based on the secrecy of a key
 - It is easier to keep a secret key than a secret algorithm
 - Keys can be more easily changed than algorithms
 - Communication with many parties is easier
 - Published designs undergo public review
- *“A cryptographic scheme for a given task is secure, if no adversary of a specified power can achieve a specified break”*

Cryptographic Devices


- are electronic devices that implement cryptographic algorithms and that store cryptographic keys [DPAbook]
- their security should not rely on the secrecy of its implementation
 - *no security by obscurity*
- Breaking a cryptographic device means extracting the key of the device not breaking the underlying algorithm.

Cryptographic Devices, Examples

- Pay TV system
 - The broadcasting company gives smart cards to their customers and the customers pay for viewing the content.



- The broadcasting company cannot assume that all customers are honest. In fact, a customer might want to duplicate her card to enable also her friends to watch the content



Cryptographic Devices, Examples (cont'd)

- Electronic Money (inside the chip card of the bank cards)
- ID cards, passports
- Car Remote Controls/Immobilizers
- USB Tokens
- RFID tags
 - (smart cards were the starting point,
but still are sensitive cases)
- Several cryptographic devices are involved in our daily life
 - There are also a huge interest to break them
 - For financial gains, fun, illegal copying

Attack Scenarios and Goals

- The attacker is a legitimate owner of the device
 - She can use the device showing that she is an honest owner
 - She can manipulate, disturb, and observe the device
- The goal of the attack is to find the secret key
- Once the secret key is revealed, the security is gone....
- Later she may use the revealed key for other purposes
 - Deciphering the communication of the others
 - Counterfeiting the contents, ...
 - Illegal access to the secured objects

Examples of Attacks on Crypto Devices

- Physical attack means to use the physical properties of the cryptographic device in order to break it.
 - then, PHYSICAL ACCESS to the device is needed
- We start with an example:
 - Eve finds a USB stick with secret data on it
 - In order to access this secret information an 8-digit PIN needs to be entered correctly
 - The device has a delayed response - so one trial takes 1s.
 - For a Brute-Force Attack, there are 10^8 possible PIN codes
 - Trying all possible PINs takes $10^8 \text{ s} \approx 1157 \text{ days} \approx 3 \text{ years}$
- Can she find a better attack?

Examples of Attacks on Crypto Devices

- Suppose that she knows the implementation platform
 - a processor
 - A part of the software is responsible for PIN checking
 - The checking is done digit by digit (like strcmp)
 - Once the wrong digit is detected, the function terminates
 - Execution time of this process is data-dependent

How to attack?

Attacking strcmp

- She should be able to measure the time
 - from sending the PIN (pressing the ENTER) to get the response
- She selects a PIN and change only the first digit while measuring the time
 - The longest time shows the correct first digit
- Finding the first one, she can do the same for the rest
- This procedure takes $10 \times 8 = 80$ timing measurements!
 - In the case that she cannot measure the timing accurately, she needs to get more timing measurements, and get the average

Attacking strcmp (cont'd)

- Running time of strcmp MUST be fixed independent of the given PIN and the stored one
 - is not easy and straightforward, but feasible

Does Eve have any further chance?

YES!

- Most of the digital circuits are implemented by CMOS technology
 - Power consumption of the processor depends on the **executed instruction** and **processed data**
 - We will see them in more detail

Attacking strcmp (cont'd)

- She should be able to measure the power consumption during PIN checking process

[The Attack scenario is similar to the last one]

- She selects a PIN and change only the first digit while measuring the power consumption
 - The most different power consumption shows the correct first digit
- Finding the first one, she can do the same for the rest
- Complexity and noise problem, the same as timing attack

Attacking strcmp (cont'd)

- Overcoming this problem is not easy
 - The power consumption of the device should be somehow independent of the processed data/the executed path
 - Or the interesting part for Eve should be hidden somewhere
 - We will see the countermeasure scheme in detail

Does Eve have any further chance?

YES!

Attacking strcmp (cont'd)

- Digital circuits need special conditions to operate correctly
 - Frequency of operation
 - Temperature
 - Voltage
- If Eve can change this condition while entering the PIN
 - by changing the frequency (clock glitch)
 - by changing the supply voltage (power glitch)
 - She may be able to skip the instruction checking the PIN
 - Yes, is hard and needs deep knowledge about the design/when the PIN is checked/...
 - If so, she does not get the PIN, she can use it once and pass the authentication

Physical Attacks

are classified into groups:

- From the attacker's behavior point of view
 - Passive Attacks
 - The device works normally; the attacker just listens
 - Active Attacks
 - The attacker makes trouble for the device
- From invasiveness point of view
 - Invasive Attacks
 - Without any limit; the strongest attacker model
 - Semi-Invasive Attacks
 - The device is depacked, but no electrical contact
 - to read/induce faults
 - Non-Invasive Attacks
 - No modification at all

Physical Attacks (examples)

	Passive	Active
Invasive	Probing	Forcing Permanent Faults
Semi-Invasive	Optical Inspection	Light Attack Radiation Attack
Non-Invasive	Side-Channel Attacks	Clock Glitch Power Glitch Temperature

Side-Channel Attacks: **Power Analysis/EM Analysis/Timing Analysis**

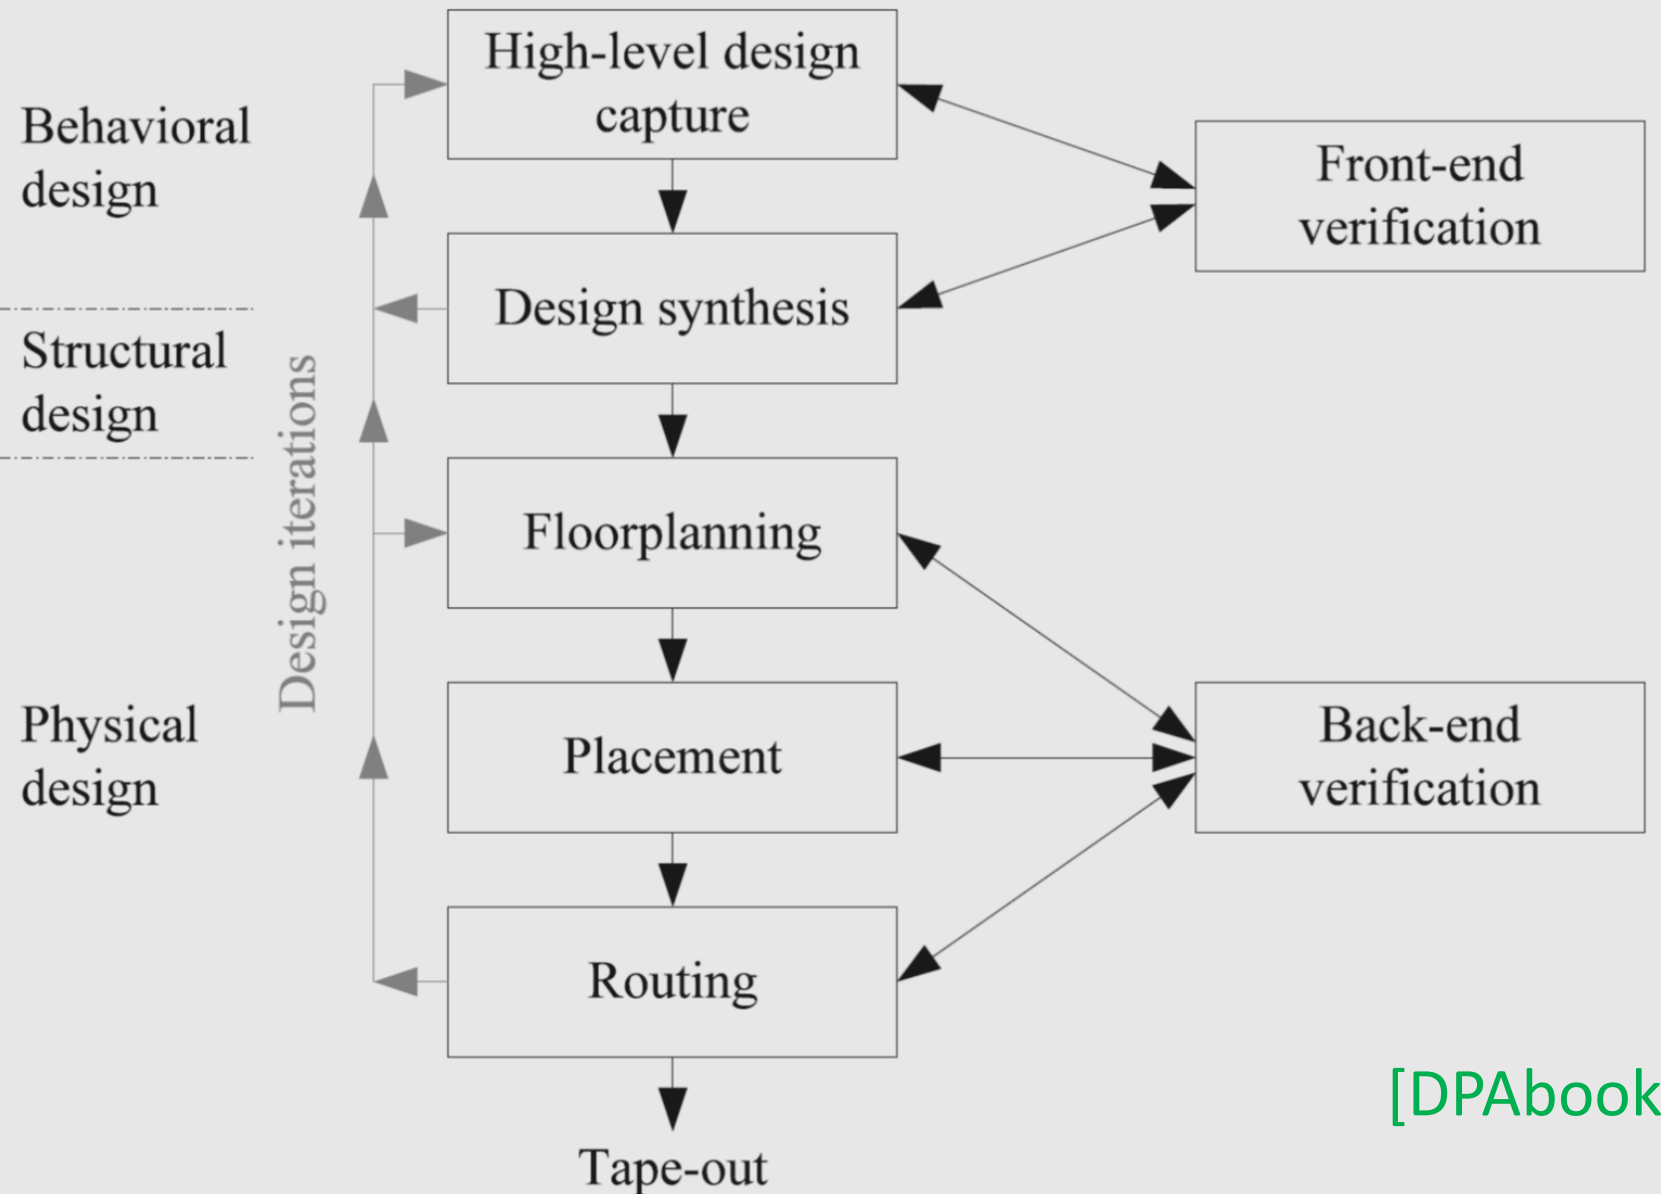
Basics of Hardware Design

- Traditionally divided into (iteratively)
 - Specification (High-level prog. lang.)
 - Behavioral Design (RTL)
 - sequences, bit widths, co-design, synch/asynch
 - HDL code
 - Structural Design
 - netlist transistors/cells
 - Physical Design
 - placement, routing

Semi-Custom Design


- A standard-cell library is used
- Many things can be done automatically
 - from “behavioral design” to the end
 - called SYNTHESIS
- The output of the synthesis is a netlist. In fact, it is a graph, where each node corresponds to a gate
- When making the LAYOUT, the gates are placed and connected
- The layout is a geometrical description of the design. It describes which material should be placed where

Semi-Custom Design (cont'd)



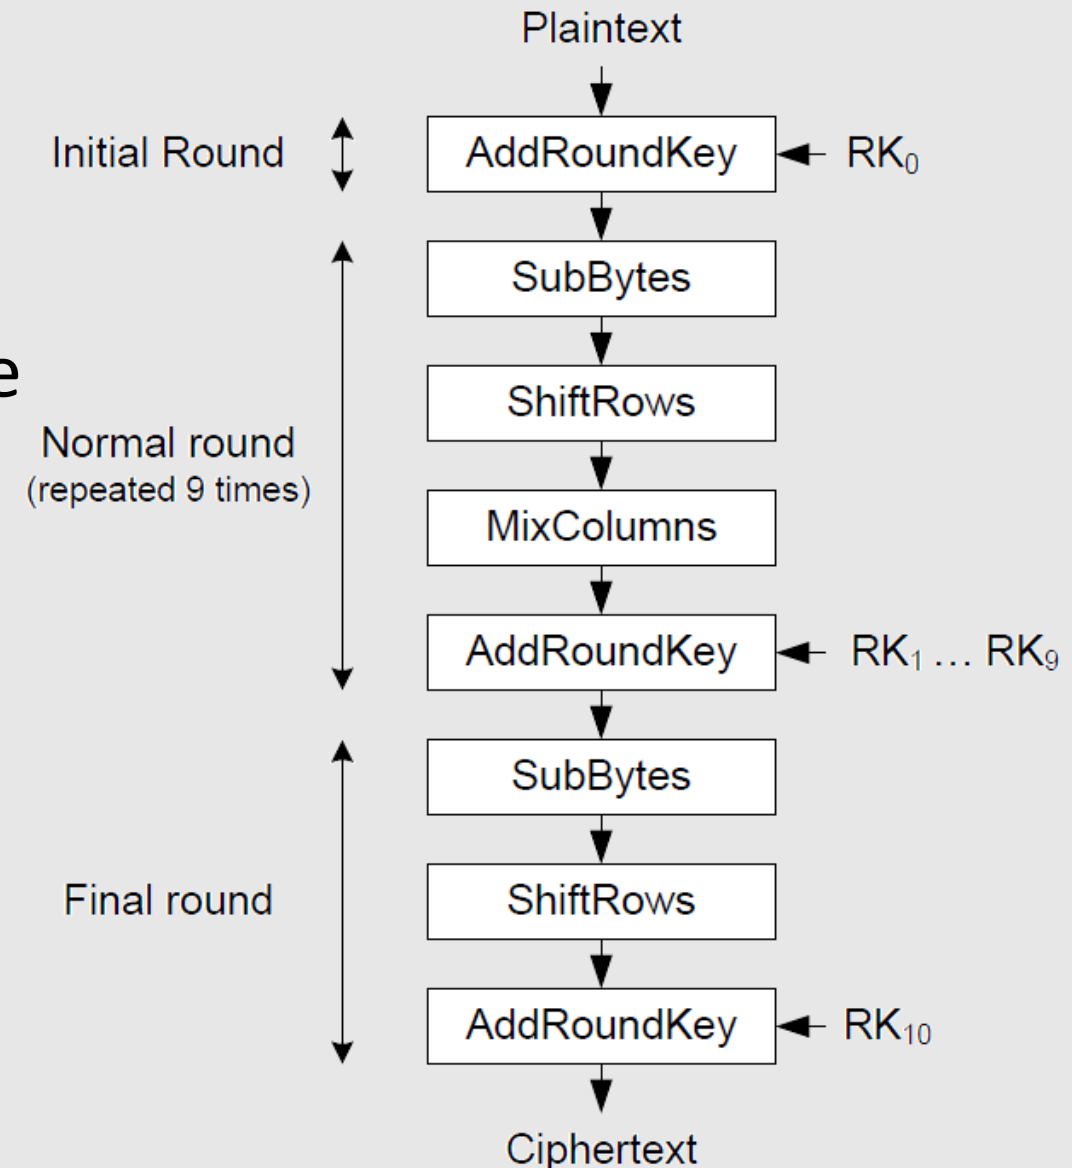
[DPAbook]

An Example




Short Review of AES

- There are 3 different versions: 128-, 192-, and 256-bit key size
- We consider only the 128-bit version.
- Plaintext and ciphertext are always 128-bit wide
- The output is ready after 10 rounds of computations
 - the last round is different



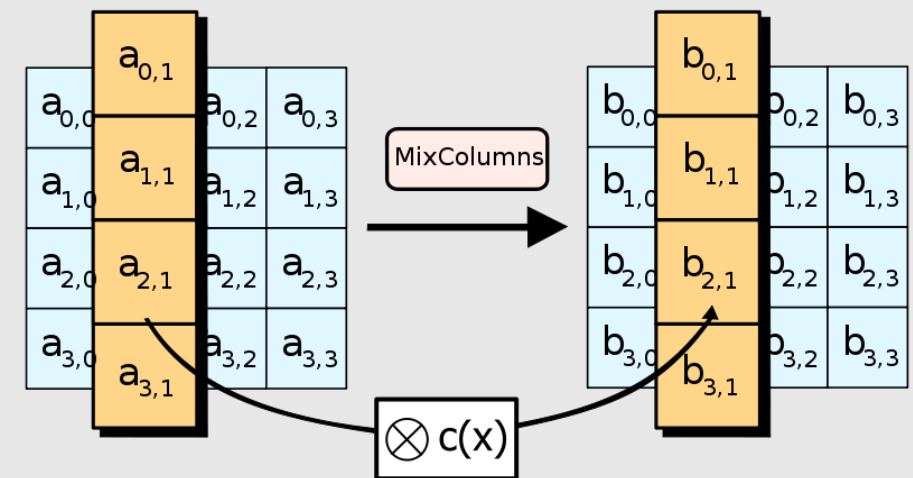
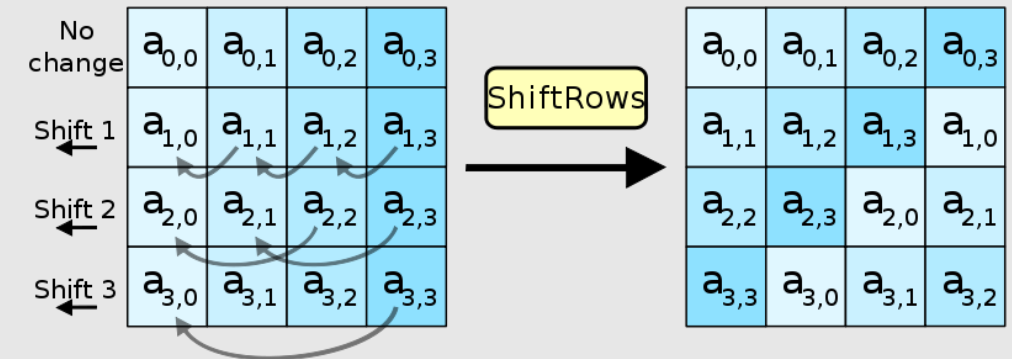
AES (AddRoundKey and SubBytes)

- All the computations are done on a 4×4 byte-wise matrix
 - All the fundamental operations are defined considering matrix representation
 - AddRoundKey is a bitwise XOR operation
 - runs over two 128-bit matrices
 - SubBytes runs the Sbox over each byte independently
 - all the Sboxes are the same
 - is an inversion in $GF(2^8)$
- in addition to an affine transformation



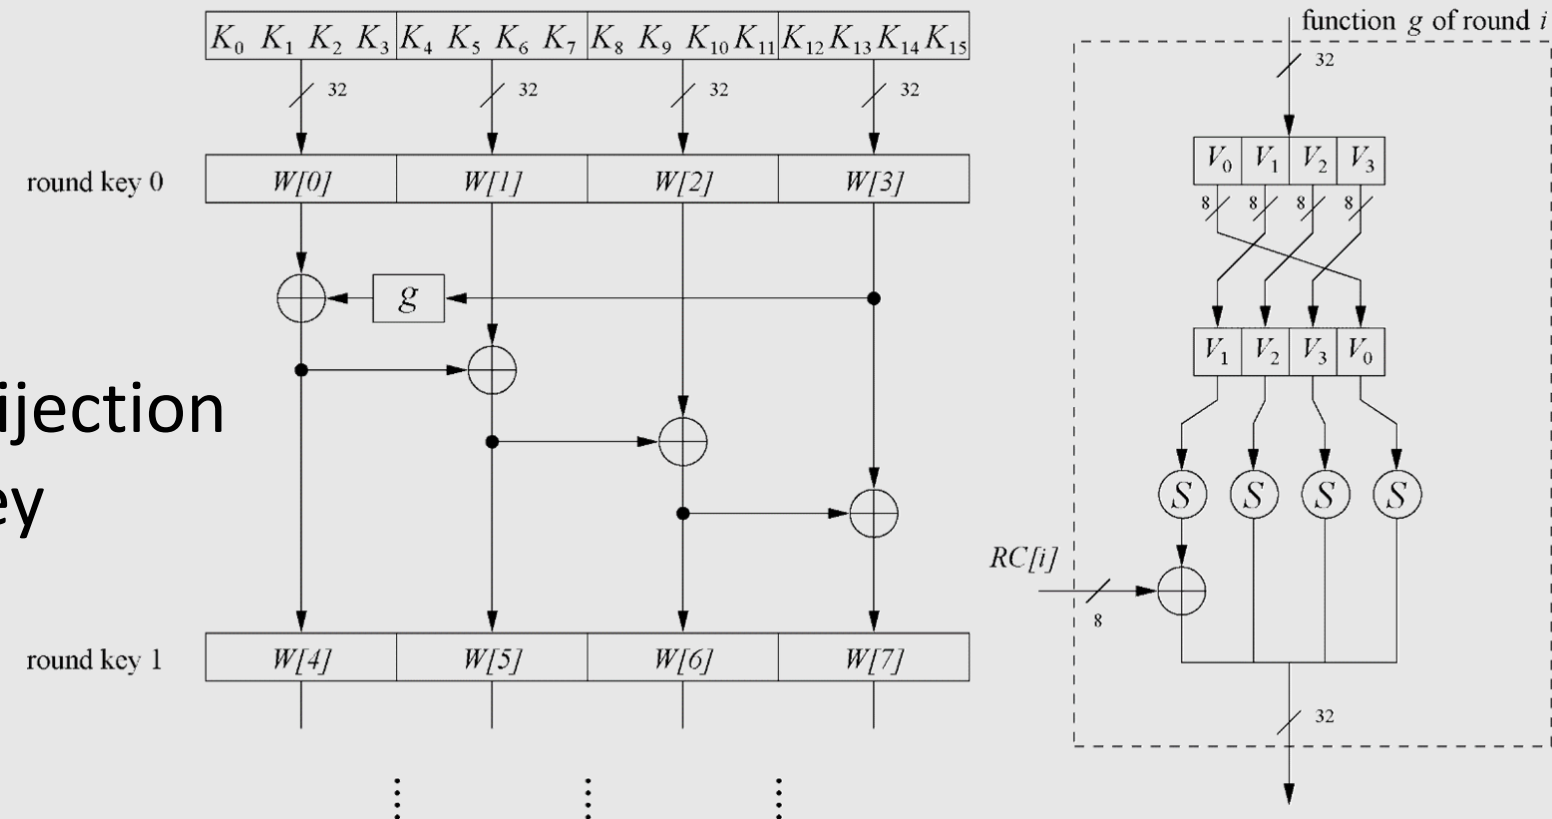
AES (ShiftRows and MixColumns)

- ShiftRows rotates 1st row, 2nd row, 3rd row, and 4th row 0, 1, 2, 3 bytes to the left respectively
- MixColumns runs on each column separately
- Each column is multiplied to a fixed matrix
 - the multiplication is by means of the same polynomial as in Sbox



AES (key schedule)

- gets 128-bit main key and provides ten 128-bit roundkeys
 - 11 roundkeys are needed, the first one is the main key
- A byte-wise rotation, 4 Sboxes, and XOR operations are used to make one 128-bit roundkey from the last roundkey
- RC[] is a round constant and is
1, 2, 4, 8, 10, 20, 40, 80, 1B, 36
(all in HEX)
- Each round of key schedule is a bijection
knowing a round key, the main key
can be trivially computed



AES (how to implement)

- Both hardware and software implementations are straightforward
- Sboxes are too big for the hardware case unless implemented by lookup tables (RAM/ROM cells)
- MixColumns is the most expensive operation in processors if the memory is restricted
- 4 Sboxes and a MixColumns can be merged to make an 8-bit to 32-bit lookup table known as T-table
- When enough memory (4kb) is available, both software and hardware implementations contain only the T-tables and XORs, which is the fastest way to realize AES

Pseudo-code AES

INPUT: Plaintext p

OUTPUT: Ciphertext c

$x = p$

$x = x \oplus k_0$

for $i = 1$ to 10

$x = \text{SB}(x)$

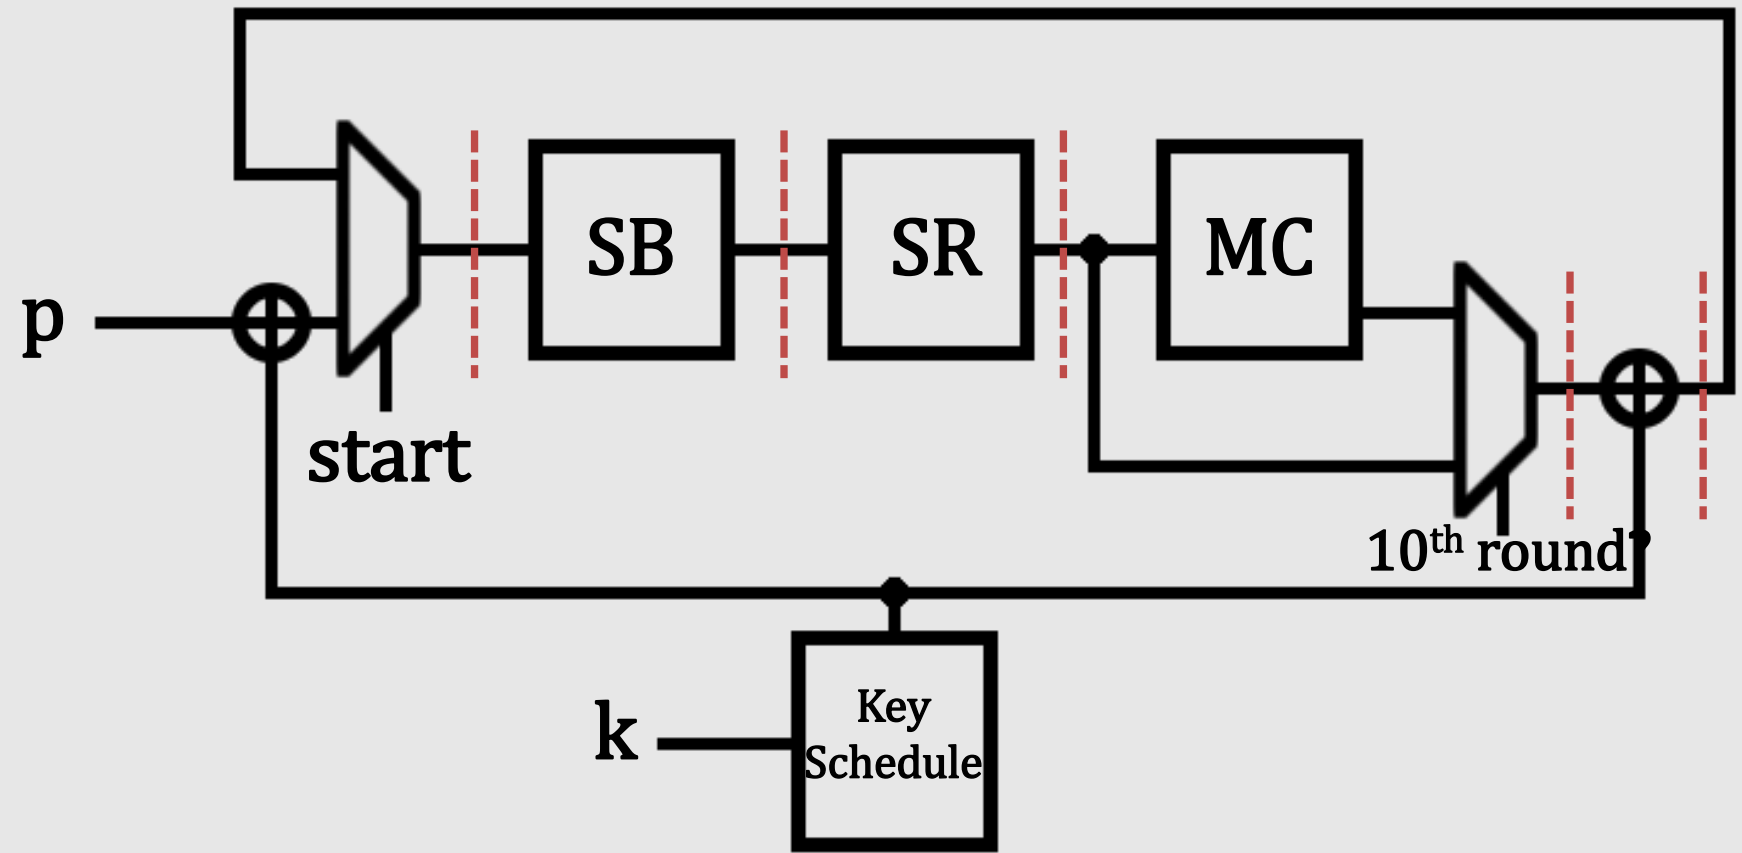
$x = \text{SR}(x)$

 if $i \neq 10$

$x = \text{MC}(x)$

$x = x \oplus k_i$

$c = x$



- It does not work without registers
- Registers can be placed between any two stages

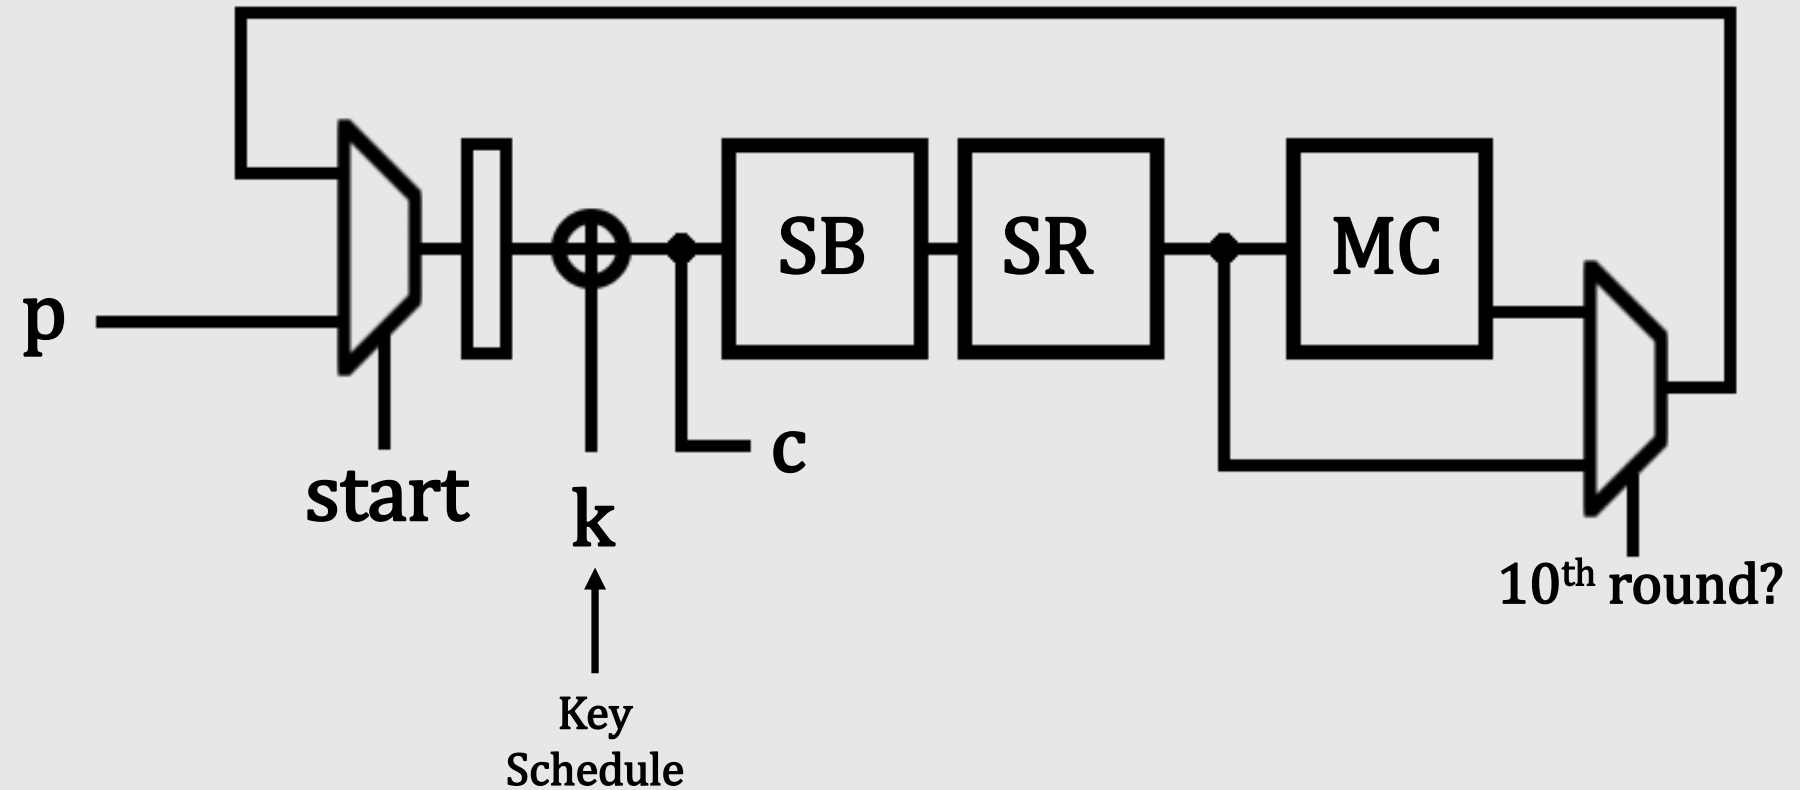
Pseudo-code AES (optimized)

INPUT: Plaintext p

OUTPUT: Ciphertext c

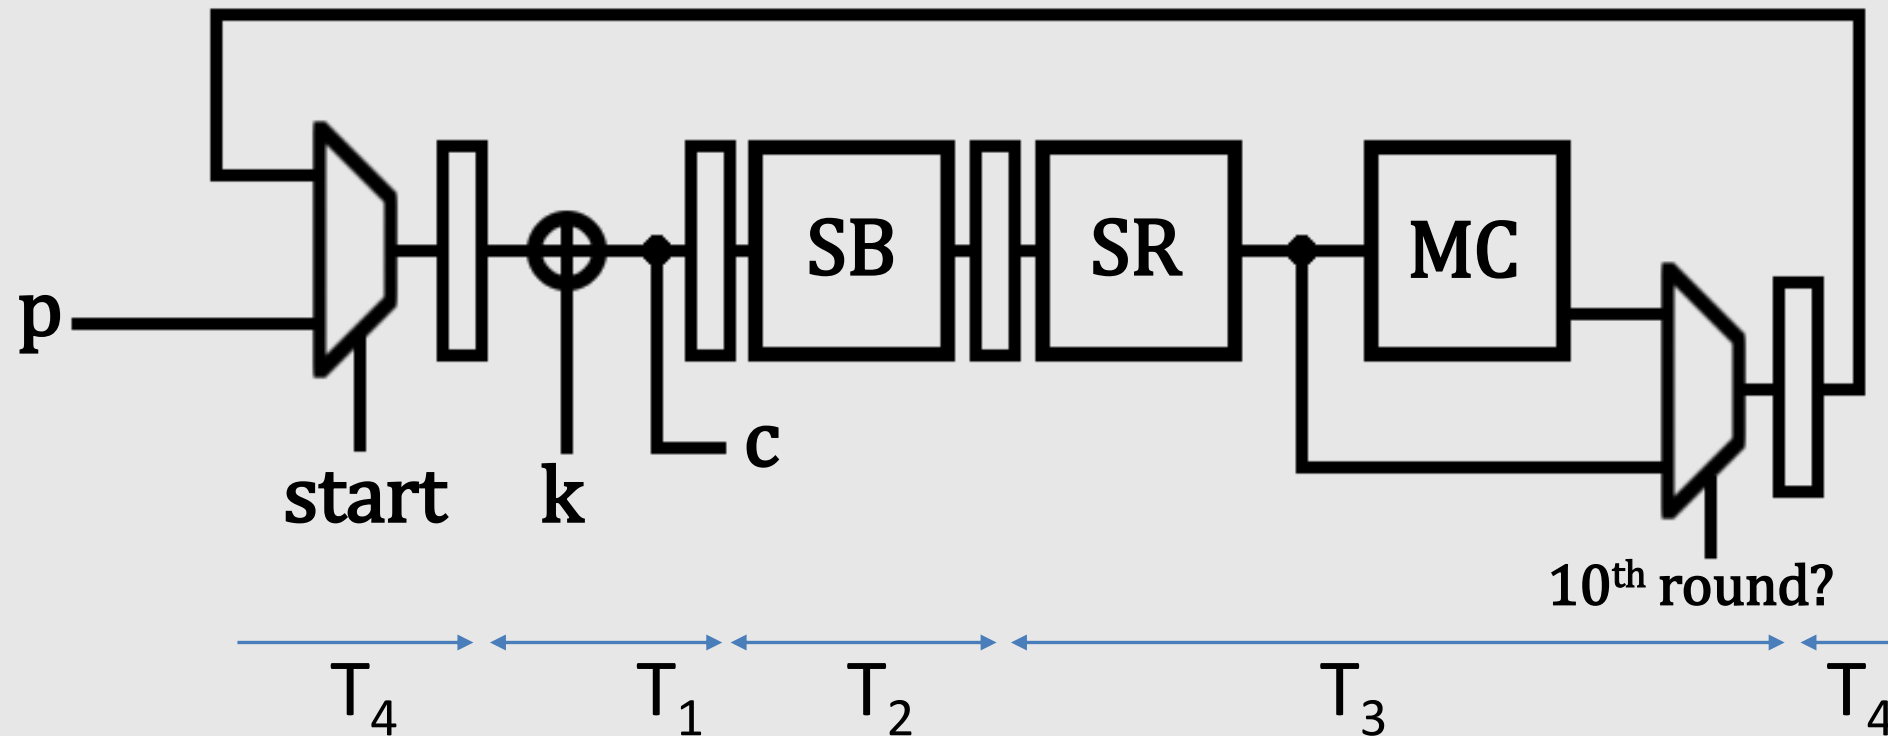
```

x = p
for i = 0 to 10
    x = x ⊕ ki
    if i ≠ 10
        x = SB(x)
        x = SR(x)
        if i ≠ 9
            x = MC(x)
c = x
    
```



Critical Path Delay

The maximum clock frequency depends on the critical path delay: the longest path between two consecutive registers.



$$\text{Max. } f = \frac{1}{T} = \frac{1}{\max(T_1, T_2, T_3, T_4)}$$

Maximum Clock Frequency


- Why ShiftRows and MixColumns are packed together?
 - ShiftRows in hardware can be realized without any combinatorial component, and has (in an ideal model) no delay.

0	4	8	12
1	5	9	13
2	6	10	14
3	7	11	15

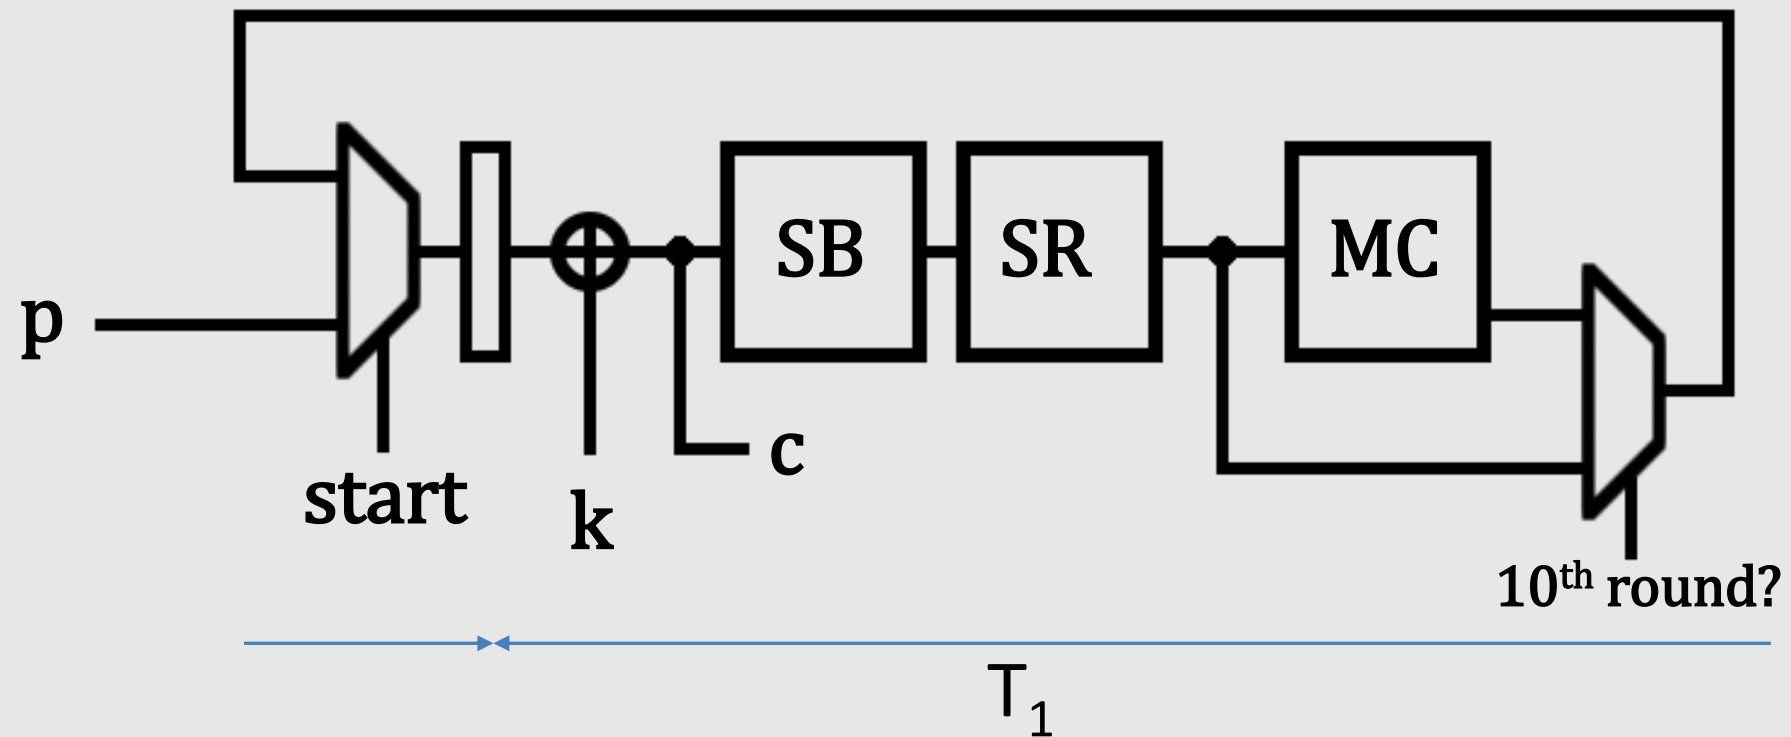


0	4	8	12
5	9	13	1
10	14	2	6
15	3	7	11

Shift Rows




Critical Path Delay



$$\text{Max. } f = \frac{1}{T} = \frac{1}{T_1}$$


Unrolling

- Fully combinatorial
- Fastest implementation (not WRT clock, but WRT latency)
- Very large area and long delay due to the combinatorial circuits



Pipeline Unrolling

- Insert registers in the unrolled structure
- Critical path delay becomes smaller
- Frequency can be increased



Pipeline Unrolling

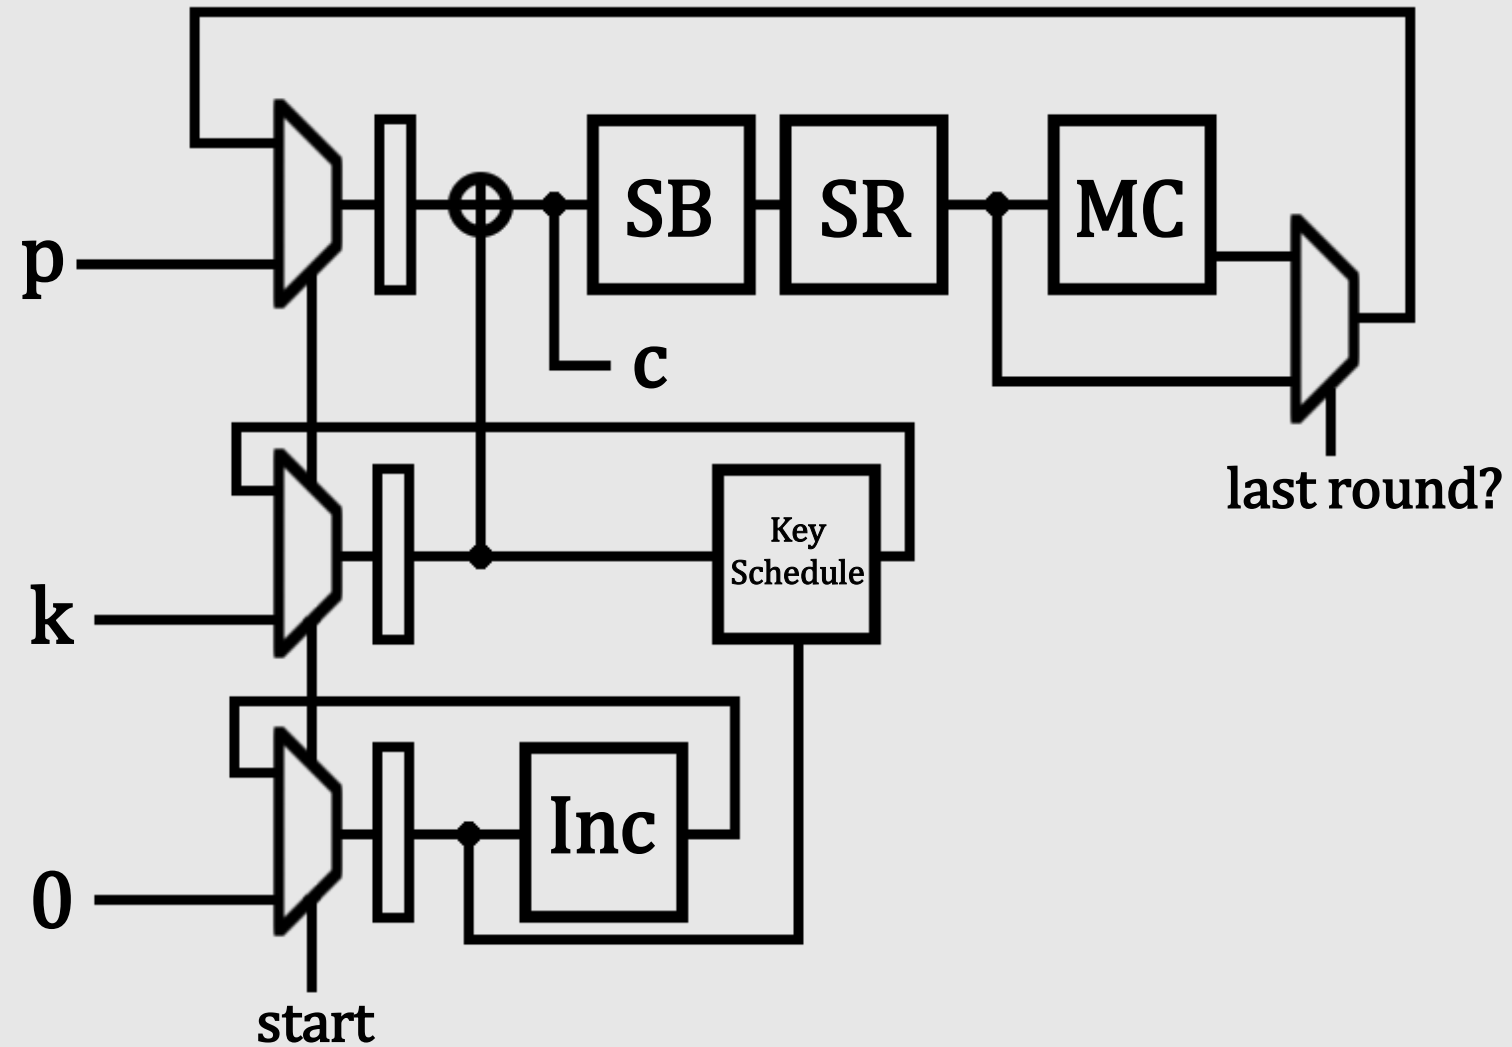
Consider an architecture with 10 registers:

- Latency of 10 clock cycles
After 10 clock cycles is the first ciphertext on output
- THEN: at each clock cycle, a further ciphertext is given as output

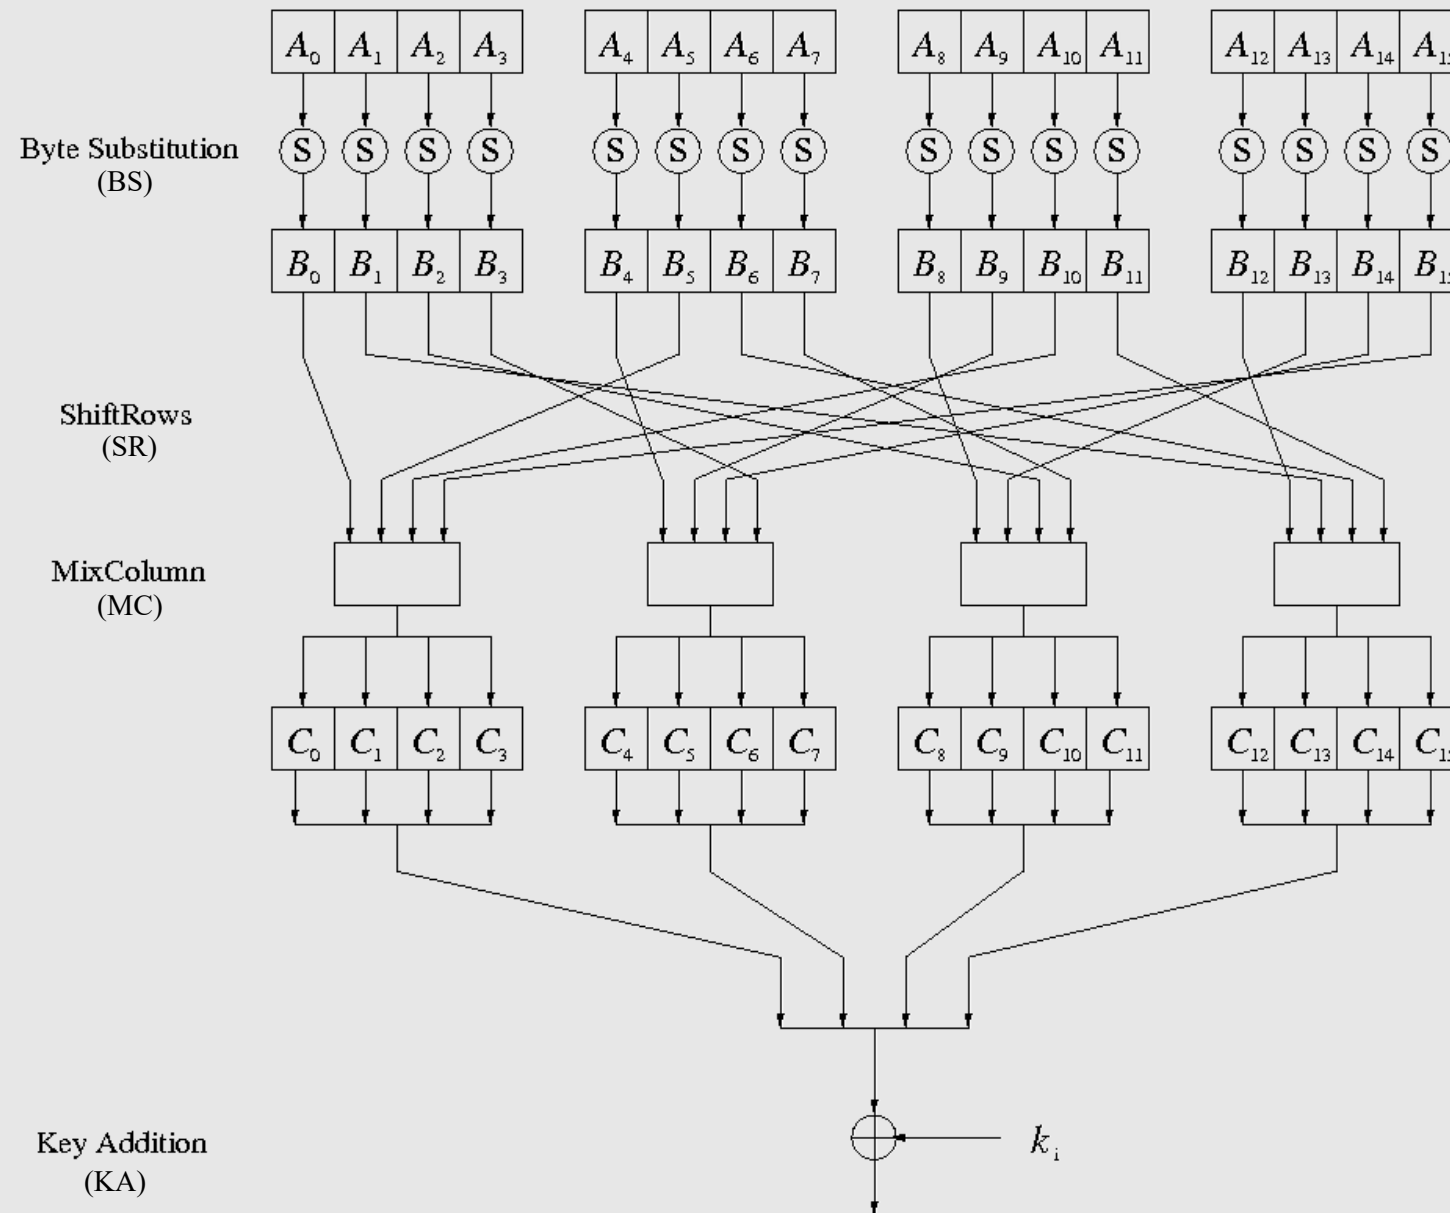
In practice, it is not usually used, but good for high throughput purposes

→ COPACOBANA (Riviera): DES with unrolled pipelining
120 FPGAs, 8 DES/FPGA,...
finding the key in a week

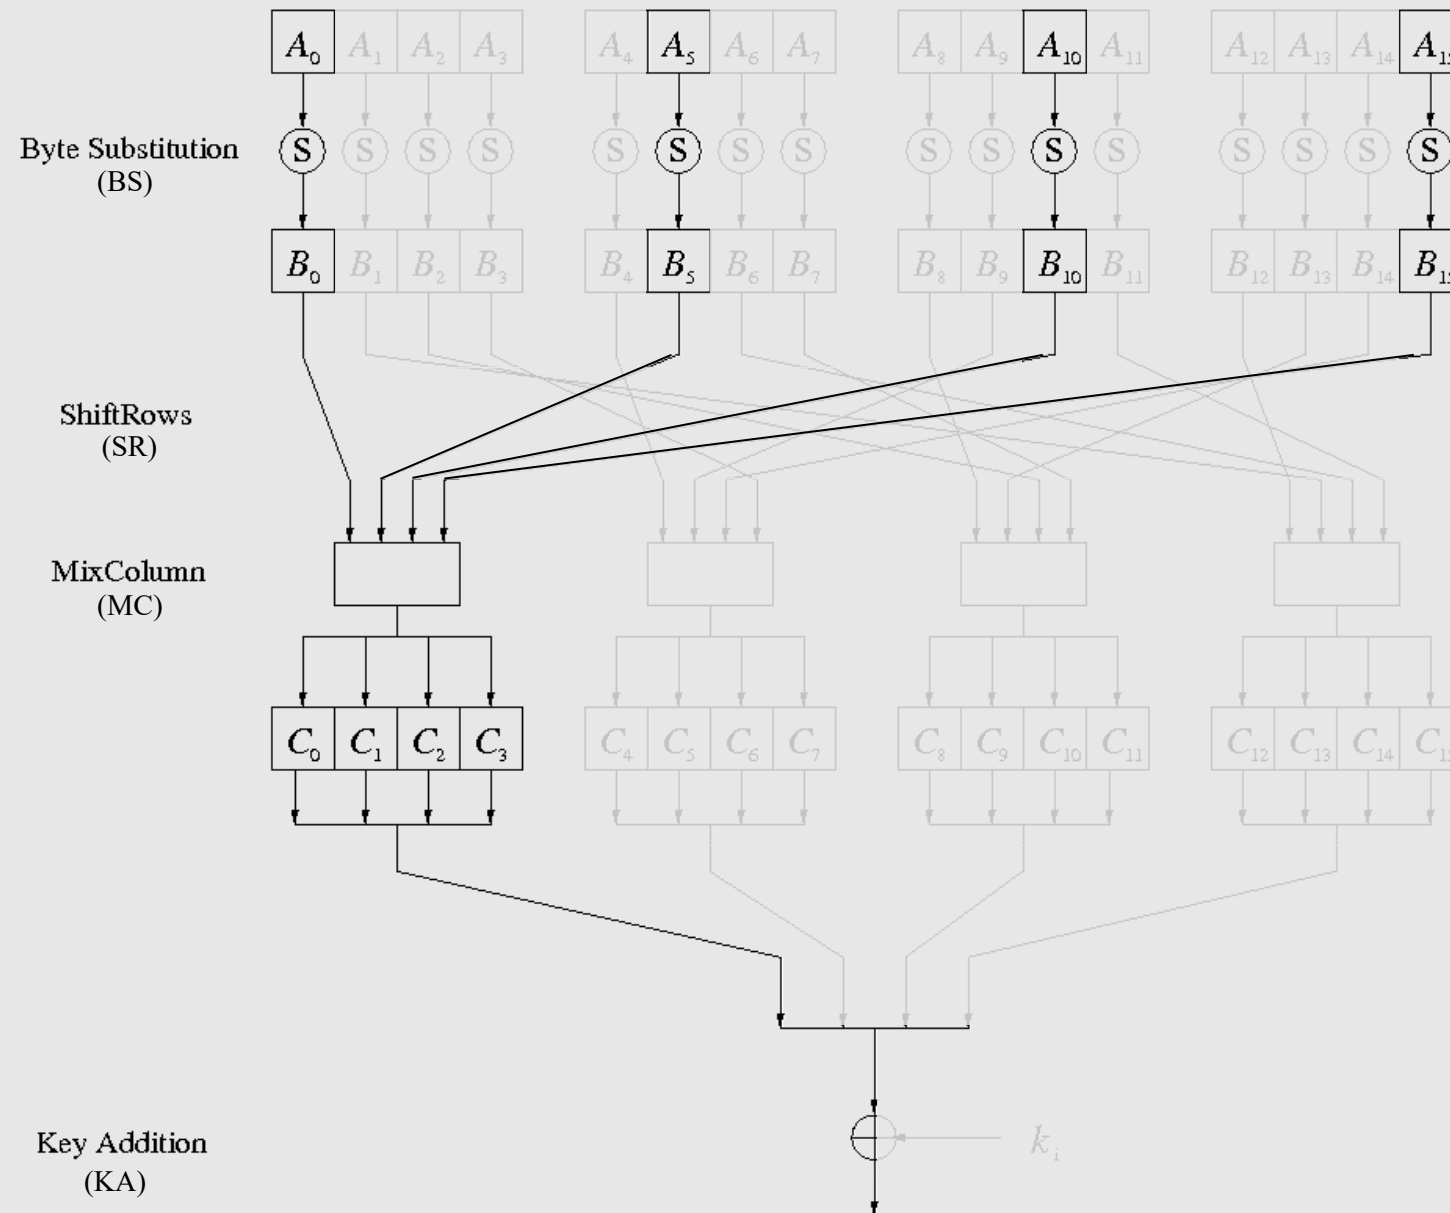
Round-Based AES



Round Function of AES-128



Round Function of AES-128



T-Table Implementation

- The most hardware designs use the 8-bit structures of the AES
- There is an alternative approach for 32-bit processors to integrate the primitive operations in the S-Box
- This results in four $T_i[x]$ tables with 256 x 32-bit entries (that is 8192 bits per table)
- These tables show the AES encryption as

$$C_j = K_{r,j} \oplus T_0[a_{4j}] \oplus T_1[a_{(4j+5) \bmod 16}] \oplus T_2[a_{(4j+10) \bmod 16}] \oplus T_3[a_{(4j+15) \bmod 16}]$$

where C_j represents 32 bits of the AES round and $K_{r,j}$ a 32-bit round key

$$T_0[x] = \begin{bmatrix} S[x] \times 02 \\ S[x] \\ S[x] \\ S[x] \times 03 \end{bmatrix}$$

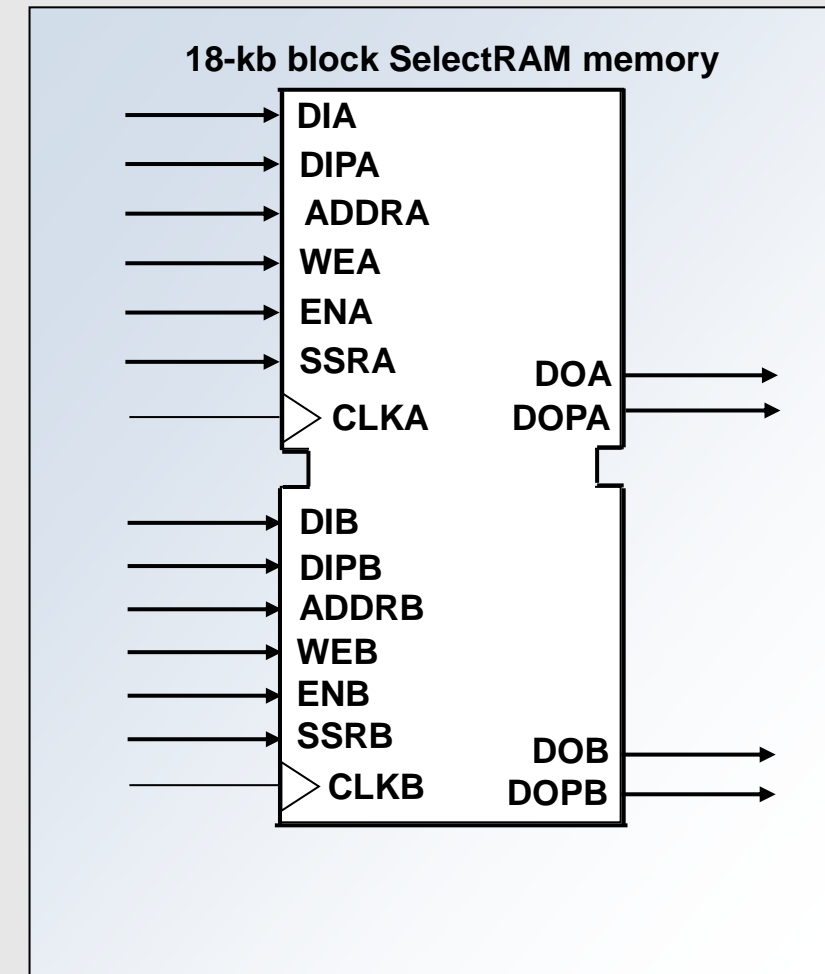
$$T_1[x] = \begin{bmatrix} S[x] \times 03 \\ S[x] \times 02 \\ S[x] \\ S[x] \end{bmatrix}$$

$$T_2[x] = \begin{bmatrix} S[x] \\ S[x] \times 03 \\ S[x] \times 02 \\ S[x] \end{bmatrix}$$

$$T_3[x] = \begin{bmatrix} S[x] \\ S[x] \\ S[x] \times 03 \\ S[x] \times 02 \end{bmatrix}$$

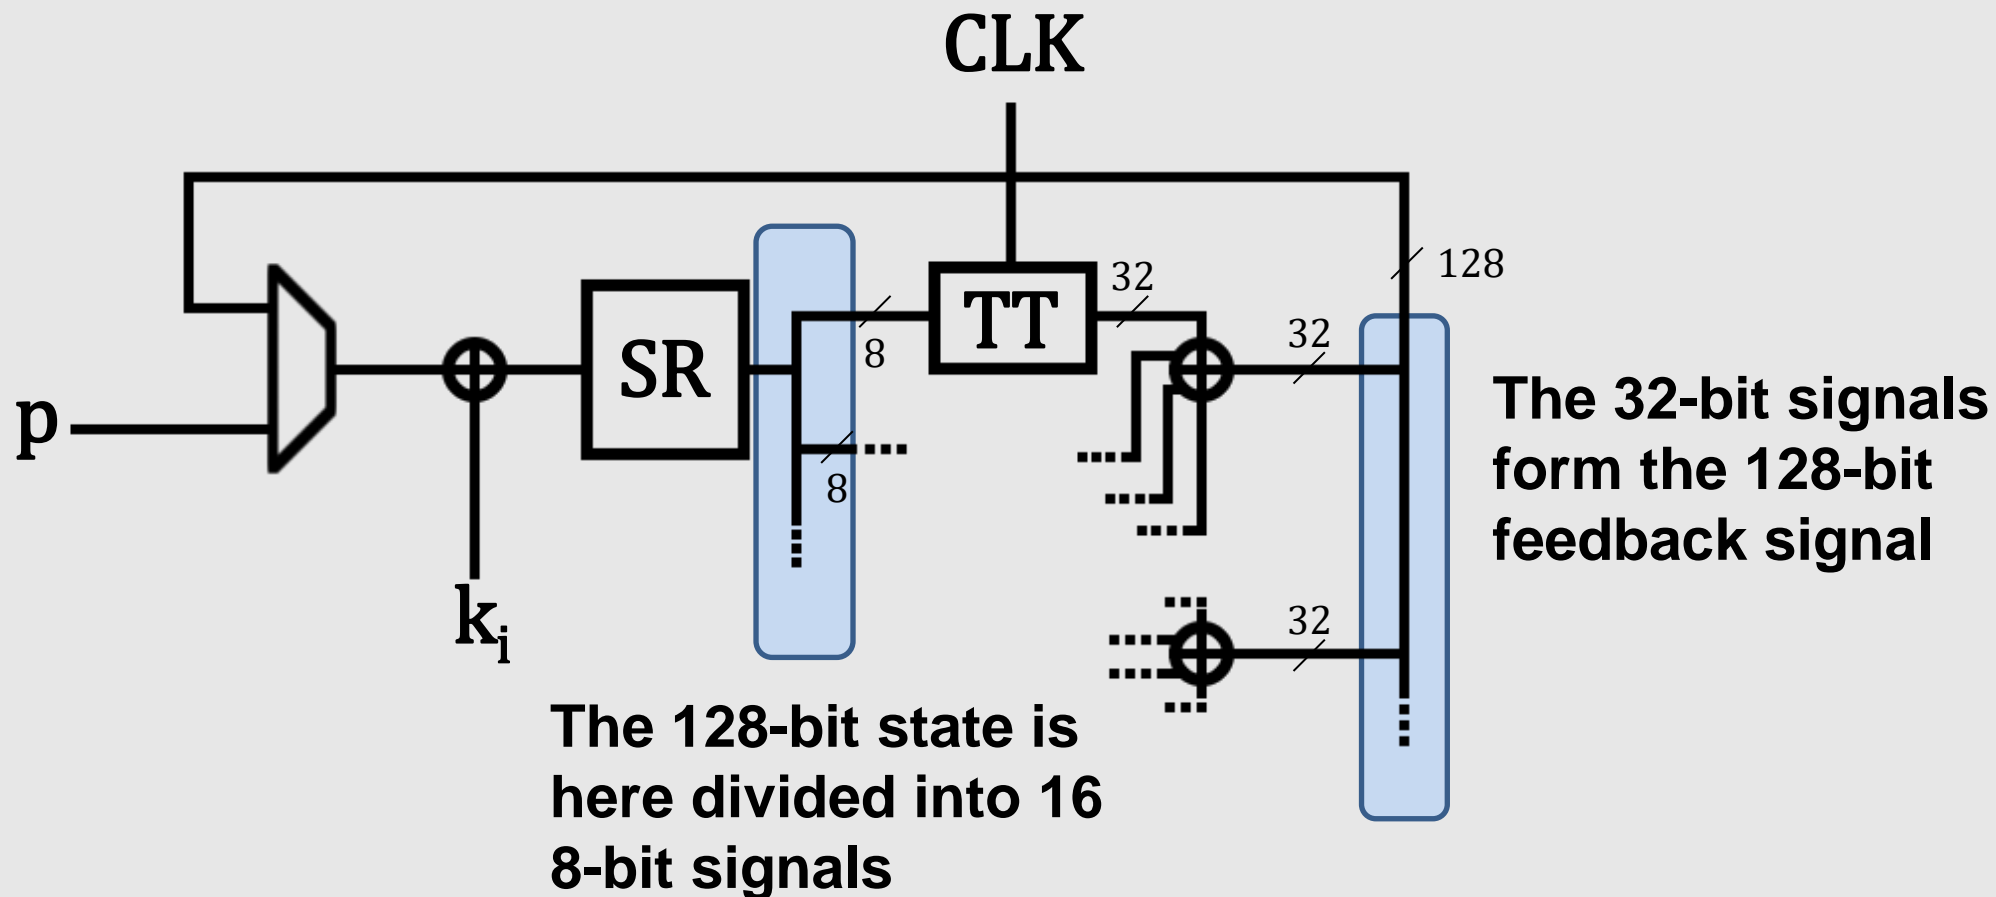
FPGA Block Memory (BRAM)

- Dedicated memory in 18/36-Kbit blocks
- True-Dual-Port
Two Ports: each one can at the same time and independently be accessed for read and write




T-Table AES (making use of FPGA BRAMs)

- BRAM is „clocked“
 - It means that no extra register is necessary.
- True Dual-Port also allows to employ only one BRAM for two T-Tables.



T-Table Last Round



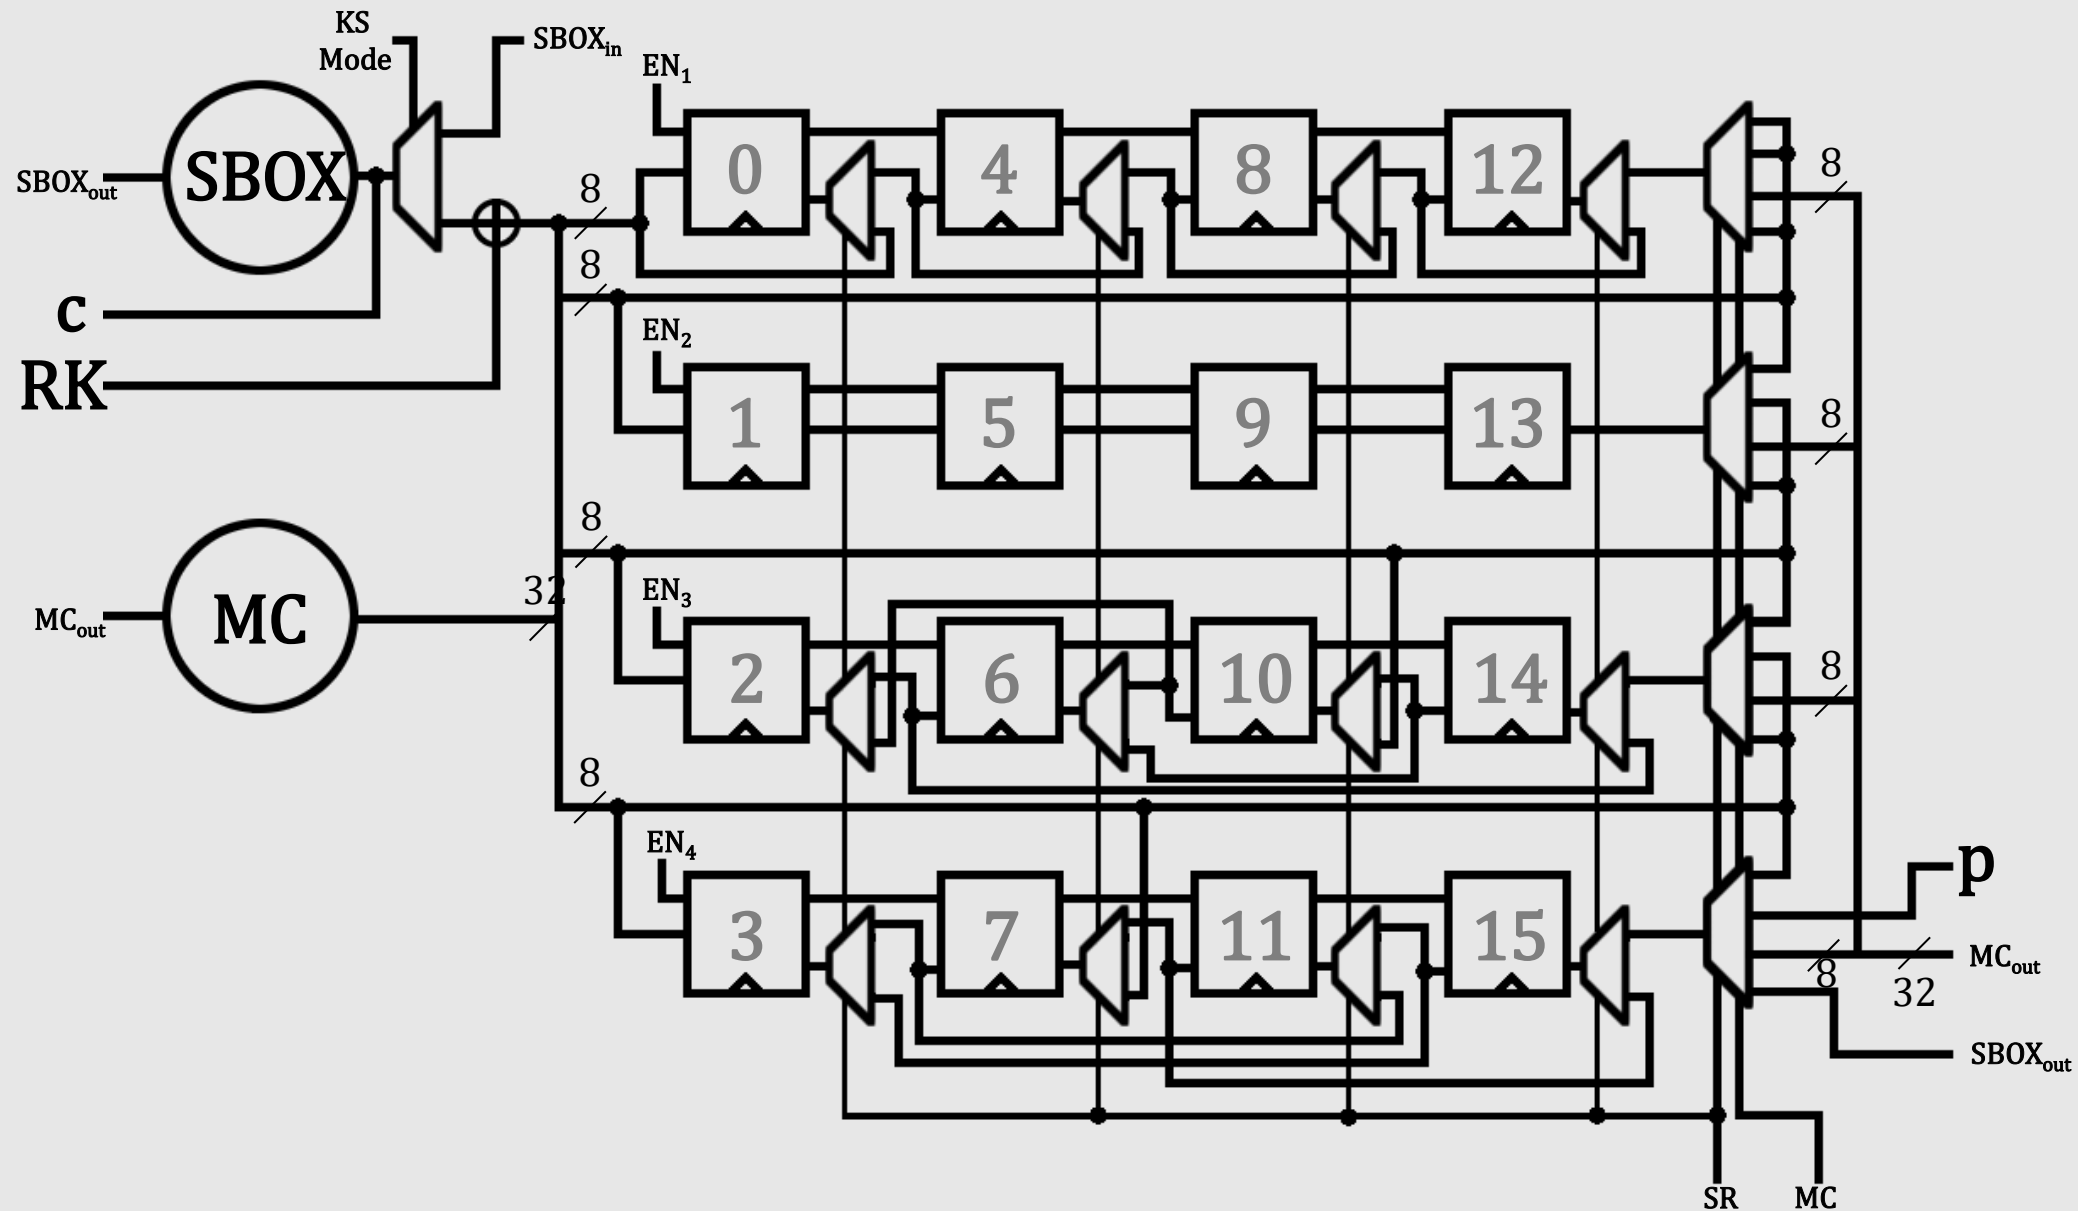
T-Table AES

- 16 T-Tables are required
- 8 BRAMs are enough
 - Because of the True Dual-Ports
- Very fast, but a large implementation wrt BRAMs

AES with Small Area

- The Sboxes need very much area
- Idea: instantiate only one Sbox module and share it
- This results in „Serialized AES“

Serialized AES with MC in one clock cycle



Serialized AES: Key Schedule

